

AETHYX LABS

AethyxLM Business Model and Launch Strategy

India-first consumer AI powered by multilingual intelligence, persistent user-owned memory and efficient context handling.

STRATEGIC POSITION

The AI that understands your language, remembers your world and finishes the work.

Working tagline

Tumhari language. Tumhara context. Tumhara AI.

Prepared September 2026 | Target public launch: before the end of 2027

Internal planning document. Financial estimates are directional and require validation before capital commitments.

01 Executive summary

AethyxLM should launch as a freemium India-first personal AI product, not initially as a company selling model weights or generic API access. Free chat acquires users; subscriptions monetize persistent memory, larger workspaces, higher usage and agentic actions. The context-compression system remains proprietary infrastructure and is marketed through consumer outcomes rather than technical claims.

The first beachhead is Indian students, job seekers and early-career professionals who regularly mix English with Indian languages. Aethyx must become unusually useful for their continuing work: learning across a semester, managing a job search, understanding documents and carrying projects forward without repeated explanations.

The recommended commercial structure is:

- Free chat with controlled usage and limited memory.
- Aethyx Plus at approximately Rs 249 per month for persistent memory, document spaces, voice and higher limits.
- Aethyx Pro at approximately Rs 599 per month for advanced reasoning, multiple agents and larger limits.
- Separate agent-credit packs for costly multi-step actions.
- Partnership revenue later; advertising should remain a small, free-tier-only revenue stream.

A serious independently trained consumer model is not compatible with the current Rs 3-4 lakh budget. The recommended capital target is Rs 50-75 crore, with Rs 75-100 crore providing safer coverage for model development and two years of operations. Until then, Aethyx can validate product demand with smaller internal models and controlled third-party foundations, without misrepresenting model provenance.

Decision	Recommendation
Company category	Consumer AI application company with proprietary model and context technology.
Launch geography	India first, beginning with concentrated student and early-career communities.
Core promise	Natural Indian communication, continuity and useful action.
Primary revenue	Monthly and annual subscriptions.
Technical target	A deeply trained 8B-12B dense model before pursuing larger parameter counts.
Defensible assets	Indian interaction data, context compiler, post-training stack, agent workflows, trust and distribution.

02 Market reality and strategic premise

The opportunity is large: India had approximately 1.03 billion internet subscriptions reported for December 2025, and the market is overwhelmingly mobile. This makes Android performance, low bandwidth, voice and affordable pricing central product requirements.

However, the competitive standard is already extremely high:

- Gemini supports more than 25 Indian languages and dialects in voice and text.
- Gemini offers personal intelligence connected to Google services.
- ChatGPT Go already links affordable pricing with longer memory, projects and expanded usage.
- Google has used large free student promotions in India, making acquisition through price alone difficult.

Therefore, 'Indian languages', 'memory' and 'agents' are not individually defensible claims. Aethyx must deliver a coherent experience that competitors do not prioritize: culturally natural code-switching, transparent user-owned memory, local workflows, low-bandwidth operation and dependable continuity across months of work.

Strategic principle

Aethyx should not try to be better than every frontier assistant at every task. It should be close enough on everyday intelligence while decisively better for a narrow set of Indian consumer journeys.

Launch condition

Marketing should not begin at scale until blind evaluations show strong everyday quality, clearly superior performance on selected multilingual workflows, reliable memory retrieval, acceptable hallucination rates and fast mobile response times.

03 Target customer and jobs to be done

Segment	Immediate needs	Reason to pay
Students	Learning, exam preparation, notes, explanations and project help.	One persistent study workspace that remembers subjects and progress.
Job seekers	Resume, applications, interview practice and skill development.	A career companion that tracks applications and continuously improves preparation.
Early-career professionals	Writing, research, coding, planning and document work.	Continuity across projects plus faster agentic execution.
Multilingual households	Translation, explanations, forms and daily assistance.	Natural speech and code-switching with less technical friction.

Primary launch persona

An 18-30 year-old Indian user on Android who switches naturally between English and an Indian language, uses AI several times each week and has recurring learning, career or project context that is painful to explain repeatedly.

Three launch workflows

- Semester companion: remembers syllabus, weak topics, uploaded notes and learning progress.
- Career companion: remembers the user's background, target roles, resumes, applications and interview performance.
- Project companion: keeps documents, decisions, tasks and preferences together across sessions and devices.

04 Product proposition and differentiation

Consumer-facing promise

Stop explaining yourself again.

Capability	Consumer benefit	What must be true
Natural Indian communication	Comfortable Hinglish, mixed-script and regional-language conversation.	Native evaluation and preference data, not translated English benchmarks.
Aethyx Memory	Projects continue without repeatedly pasting background information.	Accurate retrieval, consent controls and visible correction/deletion.
Personal knowledge spaces	Documents, chats and decisions remain organized by goal.	Reliable ingestion, citations and permission boundaries.
Local context compiler	Potentially greater privacy, lower bandwidth and longer usable history.	Compression benefits must be demonstrated experimentally before marketing.
Indian workflow agents	The assistant can complete useful multi-step work.	Every consequential action must be previewed, authorized and reversible where possible.
Efficient inference	Fast responses and affordable subscriptions.	Quantization, caching, routing and compression must lower real cost per retained user.

What not to claim

- Do not advertise a 70 percent context reduction until independently reproducible benchmarks support it.
- Do not call an adapted third-party foundation fully Aethyx-trained.
- Do not lead with parameter count. Users buy reliability, speed and outcomes.
- Do not claim privacy while targeting advertisements from private chat or memory data.

05 Business model

Plan	Indicative price	Included value
Free	Rs 0	Everyday chat, selected languages, limited messages, short memory and basic uploads.
Aethyx Plus	Rs 249/month	Persistent Memory Vault, document spaces, voice, higher limits and priority availability.
Aethyx Pro	Rs 599/month	Advanced model access, larger limits, multiple agents and deeper research/coding workflows.
Student annual	Rs 1,499-1,999	Discounted Plus access with verification and strict fair-use limits.
Agent credits	Rs 99-499 packs	Metered browser, document and multi-step automation beyond plan allowances.

Target revenue mix after product-market fit

Source	Target share	Role
Subscriptions	70%	Predictable recurring revenue and the main business.
Agent credits	15%	Pays for expensive variable-compute actions.
Partnerships and bundles	10%	Telecom, education, device and distribution partnerships.
Advertising	5%	Optional free-tier support after meaningful scale; never the primary business.

Advertising policy

Do not launch with ads. If ads are introduced later, keep them clearly labeled and outside generated answers. Private conversations, uploaded documents and personal memory must never be used for ad targeting. Paid tiers should remain ad-free.

Unit-economics rule

Do not price subscriptions using the unverified assumption that context compression will reduce cost by 70 percent. Price against measured inference cost, free-to-paid conversion, support burden and retention. Agent actions need explicit usage limits until their cost distribution is understood.

06 Model, ownership and technology strategy

The product may temporarily use an external foundation internally to validate demand, but the long-term launch model should either be independently trained, acquired with full intellectual-property assignment or covered by a negotiated commercial licence. A model based on third-party open weights remains derived from that foundation even when fine-tuned and rebranded.

Recommended technical target

- A deeply trained 8B-12B dense transformer rather than a larger but undertrained model.
- English, Hindi, Hinglish and a deliberately selected set of Indian languages at launch.
- High-quality multilingual pretraining measured in trillions of tokens, not merely billions.
- Separate conversational, reasoning, safety and agent post-training stages.
- Quantized low-latency serving plus a smaller routing or draft model.
- The context compiler developed and benchmarked independently before architectural dependence.

Why size alone is insufficient

Consumer usefulness depends on data quality, token count, post-training, retrieval, tool execution, safety and latency. An excellent 8B model can outperform an undertrained 30B model in the workflows users actually experience.

Ownership boundary

Asset	Aethyx ownership strategy
Foundation weights	Independently train or acquire full commercial ownership for the long-term model.
Context compiler	Closed proprietary technology and the core research differentiator.
Indian post-training data	Curated, licensed, documented and versioned proprietary mixture.
Evaluation harness	Private multilingual, memory, safety and agent regression suite.
Application and agents	Closed product workflows, orchestration and user experience.
User data	User-controlled; never treated as company-owned training material without specific informed consent.

07 Funding requirement

The current Rs 3-4 lakh budget supports prototypes, data-pipeline work and small-model experiments. It does not support an independently trained consumer-grade foundation model.

Capital	Realistic outcome
Rs 3-4 lakh	Research prototype and small-model experimentation.
Rs 50 lakh-Rs 2 crore	Consumer application using external weights or APIs.
Rs 10-20 crore	Proprietary 3B-7B India-focused model with limited experimentation.
Rs 40-80 crore	Strong proprietary 7B-15B consumer programme with professional data and post-training.
Rs 100-250 crore	Broader 20B-40B or MoE programme with stronger multimodal and agent capabilities.
Rs 500 crore+	Attempt to compete broadly with major international laboratories.

Evidence anchor

Meta reports 1.46 million H100 GPU-hours and about 15 trillion pretraining tokens for Llama 3.1 8B. At current IndiaAI H100 pricing, an equivalent amount of listed raw compute is approximately Rs 17-22 crore before any subsidy. This excludes failed runs, data, staff, post-training, evaluation and production inference.

Recommended capital plan

- Absolute minimum credible raise: Rs 15-20 crore for a narrower proprietary attempt.
- Recommended model-and-product funding: Rs 50-75 crore.
- Safer funding including approximately two years of operations: Rs 75-100 crore.

Illustrative Rs 60 crore allocation

Workstream	Budget
Pretraining compute and experiments	Rs 20 crore
Data acquisition, cleaning and deduplication	Rs 8 crore
SFT, preference tuning, safety and evaluation	Rs 8 crore
Research and engineering team	Rs 10 crore
Inference infrastructure and beta usage	Rs 7 crore
Application, voice, agents and context platform	Rs 4 crore
Security, legal, operations and contingency	Rs 3 crore

08 Go-to-market and launch marketing

Campaign platform

Stop explaining yourself again.

The launch demonstration should show a user importing months of notes, conversations and documents once, speaking naturally in Hinglish, and receiving continuous assistance through an exam, job search or project without repeatedly rebuilding context.

Three-month pre-launch

- Recruit 1,000 private-beta users across 20-30 colleges and startup communities.
- Use small education, career and technology creators rather than expensive celebrity endorsements.
- Publish weekly real-workflow demonstrations and collect permissioned before-and-after stories.
- Build a waitlist whose referral system grants earlier access or Plus days.
- Run blind comparisons for Hinglish, memory and task completion; publish methodology and failures.

Public launch assets

- A 60-second hero demonstration focused on continuity rather than model specifications.
- Android-first application, seven-day Plus trial and guided Memory Vault onboarding.
- Transparent comparative benchmark and a public status page.
- Referral rewards that give both participants Plus time or agent credits.

Priority channels

Channel	Execution
YouTube	Education, career and technology creators demonstrating real workflows.
Instagram/Reels	Short multilingual comparisons and continuity demonstrations.
Campus ambassadors	Workshops, beta communities and referral acquisition.
WhatsApp sharing	Useful outputs designed to be shared without exposing private context.
Play Store	Localized listings, screenshots, reviews and strong Android performance.
Search/content	High-intent pages for interviews, resumes, learning and Indian-language assistance.

09 Messaging, content and proof

Content pillars

- Ask the same question in real Hinglish.
- Aethyx remembers my entire semester.
- I stopped repeating my project background.
- One persistent workspace for my job search.
- Regional-language demonstrations produced with native speakers.

Proof hierarchy

Level	Proof
1	Recorded real product behavior with reproducible prompts.
2	Blind user preference tests conducted with native speakers.
3	Retention and repeat-memory usage from the private beta.
4	Independent evaluation of multilingual, privacy and context claims.
5	Case studies showing measurable learning, career or productivity outcomes.

Trust rules

- Never configure competitor comparisons unfairly.
- Separate model-generated demonstrations from scripted advertising.
- Explain what memory stores and provide one-tap deletion.
- Preview consequential agent actions and require explicit confirmation.
- Publish incidents and meaningful reliability regressions.

10 Roadmap and decision gates

Period	Primary objective	Gate
Sep-Dec 2026	MVP, multilingual chat, accounts, payments and memory prototype.	Working product and baseline evaluations.
Jan-Mar 2027	Closed beta with 500-1,000 users.	Evidence of repeat weekly use.
Apr-Jun 2027	Validate retention, willingness to pay and measured context savings.	At least one segment with credible product-market pull.
Jul-Sep 2027	Secure compute/funding and train or complete the launch-model programme.	Quality and safety gates passed.
Oct-Dec 2027	India public launch with Free and Plus tiers.	Stable inference capacity and positive contribution margin path.

Core operating metrics

Metric	Initial target
Four-week user retention	At least 25 percent before scaled acquisition.

Metric	Initial target
Free-to-paid conversion	3-5 percent, tested by segment and acquisition channel.
Paid monthly churn	Below 6 percent.
Paid memory adoption	More than 40 percent of paying users use memory regularly.
Referral coefficient	Track invitations and activated users, not raw shares.
Quality	Near leading assistants on everyday tasks; materially better on selected local workflows.
Economics	Positive gross margin after inference, storage, agent and support costs.

Stop conditions

- Do not scale paid marketing when four-week retention is weak.
- Do not train a much larger model simply because the current model cannot overcome product flaws.
- Do not launch paid persistent memory until retrieval and deletion controls are trustworthy.
- Do not allow autonomous financial, legal or account-changing actions without mature safeguards.

11 Principal risks and mitigations

Risk	Why it matters	Mitigation
Incumbent feature copying	Language and memory features can be replicated.	Win through local workflows, data, cost efficiency, trust and distribution.
Model quality gap	Users will not tolerate consistently weaker answers.	Strict launch gates, routing and concentrated workflow excellence.
Underfunded pretraining	One failed large run can consume the budget.	Scaling-law experiments, checkpoints, external review and large contingency.
Privacy failure	Persistent context increases potential harm.	Data minimization, encryption, explicit consent, audit logs and deletion controls.
Agent mistakes	External actions can create financial or reputational damage.	Sandboxing, previews, confirmations, scopes and reversible execution.
Low willingness to pay	Indian consumer AI pricing is highly sensitive.	Segment by recurring value, annual plans and controlled free limits.
Advertising conflict	Behavioral ads undermine the memory/privacy promise.	Keep ads contextual, clearly separated and never based on private context.

Final strategic recommendation

Build Aethyx as a consumer product with a focused Indian wedge. Prove retention and willingness to pay before committing to a full independent-model programme. Raise enough capital to train the model properly rather than optimizing for the appearance of owning weights. The enduring company value should come from the integrated system: model, Indian data, user-owned memory, context efficiency, agents, trust and distribution.

12 Sources and assumptions

The strategy combines company assumptions supplied during planning with public information available in September 2026. Prices, product features, licences and compute availability can change. Legal, financial and licensing decisions require professional review.

1. TRAI quarterly report, Dec 2025
https://traai.gov.in/sites/default/files/2026-03/QPIR_03032026.pdf
2. OpenAI: Introducing ChatGPT Go
<https://openai.com/index/introducing-chatgpt-go/>
3. OpenAI Help: ChatGPT Go
<https://help.openai.com/en/articles/11989085-what-is-chatgpt-go>
4. Google India: Indian-language expansion
<https://blog.google/intl/en-in/company-news/deepening-our-commitment-to-indias-ai-ambition/>
5. Google India: Personal Intelligence
<https://blog.google/intl/en-in/products/personal-intelligence-launches-in-the-gemini-app-in-india/>
6. Google India: AI learning and student adoption
<https://blog.google/intl/en-in/company-news/technology/our-approach-to-ai-in-learning-and-how-students-in-india-are-embracing-the-future/>
7. Meta Llama 3.1 8B model card
<https://huggingface.co/meta-llama/Llama-3.1-8B>
8. IndiaAI compute price list
<https://compute.indiaai.gov.in/pricelist>
9. IndiaAI cloud-services policy
<https://compute.indiaai.gov.in/files/end-user-policy.pdf>
10. Training Compute-Optimal Large Language Models
<https://arxiv.org/abs/2203.15556>
11. Apache License 2.0
<https://www.apache.org/licenses/LICENSE-2.0>
12. Qwen3-8B official model card
<https://huggingface.co/Qwen/Qwen3-8B>

AETHYX LABS

Confidential working strategy - not an investment memorandum, legal opinion or audited financial forecast.